Aurora Capacity Summary (signed extract)
The application tier scales between 3 and 12 nodes.
The quorum tier is fixed at 5 nodes and is never autoscaled.
The ingest spool tier runs exactly 2 nodes in every ring.
End-to-end latency budget: 250 ms. Quorum commit budget: 40 ms.
Peak ring power draw: 42 kW per cage.
Spool retention: 36 hours at 18 MB/s peak, 2.4 TB NVMe per node.